

# Training an ML Model

---

# The Dataset

**So far, we have introduced our simulator**

The rest of our plan is as follows

- We learn an ML model
- We embed the model in a larger optimization problem
- We obtain a solution, i.e. a set of action to control the epidemics

**But which data are we going to use for training?**

# The Dataset

**So far, we have introduced our simulator**

The rest of our plan is as follows

- We learn an ML model
- We embed the model in a larger optimization problem
- We obtain a solution, i.e. a set of action to control the epidemics

**But which data are we going to use for training?**

**Since we have a simulator, we can **build** our dataset**

- This means we can generate as much data as we wish
- ...But also that we are responsible for **how to generate it**

# Building Our Dataset

**We need to define the **structure** of the dataset**

- We will focus on Non-Therapeutic Interventions (NPI)
  - E.g. mask mandates, social distancing...
- NPIs affect the  $\beta$  parameter in a SIR model
  - We will assume to have constant  $\gamma$  in our setup
- We will focus on making predictions at weekly intervals

**Therefore, we can cover our needs with...**

For the **input** part:

- The initial state ( $S, I, R$ ) and the value of  $\beta$

For **output** part:

- The state after one week ( $S, I, R$ )

**Given an input ( $S, I, R, \beta$ ), we can get the output via simulation**

# Building Our Dataset

## Which input configurations should we generate?

A training set should be representative of the test distribution

- We do not have a fixed test distribution (no test set)
- ...But we know that the ML model will be **used by an optimizer**

## The optimizer will seek to **minimize the total infections**

So, we will need:

- High accuracy on the best configurations, so as to **find** them
- High accuracy on the worst configurations, so as to **avoid** them

I.e. to be safe the model should work all across the board


## Hence, we need a method that can cover well a given input space

- The simplest approach would be use use a regular grid
- ...But that approach does not scale well

# Latin Hypercube Sampling

The method we will use is called **Latin Hypercube Sampling**

Suppose we want to sample  $m$  points for  $n$  attributes with fixed ranges




- We can view the sampling space as a hypercube
- ...Then we divide each dimension in  $n$  segments

In the example we want to sample 4 points for 2 attributes

# Latin Hypercube Sampling

The method we will use is called **Latin Hypercube Sampling**

Suppose we want to sample  $m$  points for  $n$  attributes with fixed ranges




- We sample the first point uniformly at random
- ...Then we "cover" the row and column that contain the sample

# Latin Hypercube Sampling

The method we will use is called **Latin Hypercube Sampling**

Suppose we want to sample  $m$  points for  $n$  attributes with fixed ranges



- When we take additional samples, we exclude all covered row/columns
- ...So we end up with a pattern similar to that of the figure

**LHS can cover quite uniformly a given space with relatively few samples**




# Latin Hypercube Sampling

## Let's see a practical example

Here is the result of **uniform sampling**, for reference

```
In [2]: test_nsamples, test_ranges = 12, [(0., 1.), (0., 1.)]  
X = util.sample_points(test_ranges, test_nsamples, mode='uniform', seed=42)  
util.plot_2D_samplespace(X, figsize=figsize)
```




# Latin Hypercube Sampling

Let's see a practical example

...And here is the result of classical LHS:

```
In [3]: test_nsamples, test_ranges = 12, [(0., 1.), (0., 1.)]  
X = util.sample_points(test_ranges, test_nsamples, mode='lhs', seed=42)  
util.plot_2D_samplespace(X, figsize=figsize)
```




# Latin Hypercube Sampling

The process can be further improved

E.g. after sampling we can try to maximize the minimum distance

```
In [4]: test_nsamples, test_ranges = 12, [(0., 1.), (0., 1.)]  
X = util.sample_points(test_ranges, test_nsamples, mode='max_min', seed=42)  
util.plot_2D_samplespace(X, figsize=figsize)
```



# Dataset Input

We are now ready to generate our **dataset input**

```
In [5]: n_tr, n_ts = 10000, 2000
        sir_tr_in = util.generate_SIR_input(max_samples=n_tr, mode='max_min', seed=42, normalize=True)
        sir_ts_in = util.generate_SIR_input(max_samples=n_ts, mode='max_min', seed=42, normalize=True)
        sir_tr_in.head()
```

Out[5]:

	S	I	R	beta
0	0.352572	0.512310	0.135119	0.167925
1	0.175114	0.707867	0.117019	0.100628
2	0.024010	0.836985	0.139005	0.048884
3	0.139277	0.247573	0.613150	0.367367
4	0.491191	0.195846	0.312963	0.067067

- We sample  $S, I, R, \beta$  from  $[0, 1]^3 \times [0, .4]$
- ...Then  $S, I, R$  are normalized so that their sum is 1

This will reduce in some redundancy in the dataset

# Dataset Output

We obtain the corresponding output via simulation

```
In [6]: %%time
gamma = 1/14
sir_tr_out = util.generate_SIR_output(sir_tr_in, gamma, 7)
sir_ts_out = util.generate_SIR_output(sir_ts_in, gamma, 7)
sir_tr_out.head()
```

```
CPU times: user 1.2 s, sys: 53.3 ms, total: 1.25 s
Wall time: 1.2 s
```

Out[6]:


	S	I	R
0	0.201814	0.425756	0.372430
1	0.115945	0.474359	0.409696
2	0.019150	0.511369	0.469481
3	0.078295	0.196566	0.725139
4	0.453265	0.148189	0.398546

- We picked  $\gamma = 1/14$  (this will be fixed in our use case)
- We simulate one week

# Training a Model

## We try with Linear Regression

```
In [7]: nn0 = util.build_ml_model(input_size=4, output_size=3, hidden=[], name='LR')
history0 = util.train_ml_model(nn0, sir_tr_in, sir_tr_out, verbose=0, epochs=100)
util.plot_training_history(history0, figsize=figsize)
util.print_ml_metrics(nn0, sir_tr_in, sir_tr_out, 'training')
util.print_ml_metrics(nn0, sir_ts_in, sir_ts_out, 'test')
```




Model loss: 0.0012 (training) 0.0013 (validation)  
R2: 0.95, MAE: 0.023, RMSE: 0.03 (training)

# Training a Model

## ...And with a shallow Neural Network

```
In [8]: nn1 = util.build_ml_model(input_size=4, output_size=3, hidden=[8], name='MLP')
history1 = util.train_ml_model(nn1, sir_tr_in, sir_tr_out, verbose=0, epochs=100)
util.plot_training_history(history1, figsize=figsize)
util.print_ml_metrics(nn1, sir_tr_in, sir_tr_out, 'training')
util.print_ml_metrics(nn1, sir_ts_in, sir_ts_out, 'test')
```



Model loss: 0.0002 (training) 0.0002 (validation)  
R2: 0.99, MAE: 0.0081, RMSE: 0.01 (training)

# Considerations and Next Steps

**We will save both models for later**

```
In [9]: util.save_ml_model(nn0, 'nn0')  
        util.save_ml_model(nn1, 'nn1')
```

- The network is much better in terms of accuracy
- ...But the Linear Regressor is simpler!

Hence, the approaches provide different trade offs

**We are halfway there**

We now have our ML model(s)!

- We need to understand how they can be embedded in an optimization model
- ...And we need to define our optimization model itself